

TV Script Generation

In this project, you'll generate your own [Simpsons](https://en.wikipedia.org/wiki/The_Simpsons) TV scripts using RNNs. You'll be using part of the [Simpsons dataset](https://www.kaggle.com/wcukierski/the-simpsons-by-the-data) of scripts from 27 seasons. The Neural Network you'll build will generate a new TV script for a scene at [Moe's Tavern](https://simpsonswiki.com/wiki/Moe's_Tavern).

Get the Data

The data is already provided for you. You'll be using a subset of the original dataset. It consists of only the scenes in Moe's Tavern. This doesn't include other versions of the tavern, like "Moe's Cavern", "Flaming Moe's", "Uncle Moe's Family Feed-Bag", etc..

```
In [1]: """
DON'T MODIFY ANYTHING IN THIS CELL
"""

import helper

data_dir = './data/simpsons/moes_tavern_lines.txt'
text = helper.load_data(data_dir)
# Ignore notice, since we don't use it for analysing the data
text = text[81:]
```

Explore the Data

Play around with `view_sentence_range` to view different parts of the data.

In [2]: view_sentence_range = (0, 10)

```
"""
DON'T MODIFY ANYTHING IN THIS CELL
"""

import numpy as np

print('Dataset Stats')
print('Roughly the number of unique words: {}'.format(len({word: None for word in text.split()})))
scenes = text.split('\n\n')
print('Number of scenes: {}'.format(len(scenes)))
sentence_count_scene = [scene.count('\n') for scene in scenes]
print('Average number of sentences in each scene: {}'.format(np.average(sentence_count_scene)))

sentences = [sentence for scene in scenes for sentence in scene.split('\n')]
print('Number of lines: {}'.format(len(sentences)))
word_count_sentence = [len(sentence.split()) for sentence in sentences]
print('Average number of words in each line: {}'.format(np.average(word_count_sentence)))

print()
print('The sentences {} to {}'.format(*view_sentence_range))
print('\n'.join(text.split('\n')[view_sentence_range[0]:view_sentence_range[1]]))
```

Dataset Stats

Roughly the number of unique words: 11492

Number of scenes: 262

Average number of sentences in each scene: 15.248091603053435

Number of lines: 4257

Average number of words in each line: 11.50434578341555

The sentences 0 to 10:

Moe_Szyslak: (INTO PHONE) Moe's Tavern. Where the elite meet to drink.

Bart_Simpson: Eh, yeah, hello, is Mike there? Last name, Rotch.

Moe_Szyslak: (INTO PHONE) Hold on, I'll check. (TO BARFLIES) Mike Rotch. Mike Rotch. Hey, has anybody seen Mike Rotch, lately?

Moe_Szyslak: (INTO PHONE) Listen you little puke. One of these days I'm gonna catch you, and I'm gonna carve my name on your back with an ice pick.

Moe_Szyslak: What's the matter Homer? You're not your normal effervescent self.

Homer_Simpson: I got my problems, Moe. Give me another one.

Moe_Szyslak: Homer, hey, you should not drink to forget your problems.

Barney_Gumble: Yeah, you should only drink to enhance your social skills.

Implement Preprocessing Functions

The first thing to do to any dataset is preprocessing. Implement the following preprocessing functions below:

- Lookup Table
- Tokenize Punctuation

Lookup Table

To create a word embedding, you first need to transform the words to ids. In this function, create two dictionaries:

- Dictionary to go from the words to an id, we'll call `vocab_to_int`
- Dictionary to go from the id to word, we'll call `int_to_vocab`

Return these dictionaries in the following tuple (`vocab_to_int`, `int_to_vocab`)

```
In [3]: import numpy as np
import problem_unittests as tests
from collections import Counter

def create_lookup_tables(text):
    """
    Create lookup tables for vocabulary
    :param text: The text of tv scripts split into words
    :return: A tuple of dicts (vocab_to_int, int_to_vocab)
    """
    # TODO: Implement Function
    word_counts = Counter(text)
    sorted_vocab = sorted(word_counts, key=word_counts.get, reverse=True)
    )

    int_to_vocab = {ii: word for ii, word in enumerate(sorted_vocab)}
    vocab_to_int = {word: ii for ii, word in int_to_vocab.items()}
    return vocab_to_int, int_to_vocab

"""
DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
"""
tests.test_create_lookup_tables(create_lookup_tables)
```

Tests Passed

Tokenize Punctuation

We'll be splitting the script into a word array using spaces as delimiters. However, punctuations like periods and exclamation marks make it hard for the neural network to distinguish between the word "bye" and "bye!".

Implement the function `token_lookup` to return a dict that will be used to tokenize symbols like "!" into "`||Exclamation_Mark||`". Create a dictionary for the following symbols where the symbol is the key and value is the token:

- Period (.)
- Comma (,)
- Quotation Mark (")
- Semicolon (;)
- Exclamation mark (!)
- Question mark (?)
- Left Parentheses (()
- Right Parentheses ())
- Dash (--)
- Return (\n)

This dictionary will be used to token the symbols and add the delimiter (space) around it. This separates the symbols as it's own word, making it easier for the neural network to predict on the next word. Make sure you don't use a token that could be confused as a word. Instead of using the token "dash", try using something like "`||dash||`".

```
In [5]: def token_lookup():
        """
        Generate a dict to turn punctuation into a token.
        :return: Tokenize dictionary where the key is the punctuation and the
        value is the token
        """
        # TODO: Implement Function
        tokenize_dict = {
            '.': '|period|',
            ',': '|comma|',
            '"': '|quotation_mark|',
            ';': '|semicolon|',
            '!': '|exclamation_mark|',
            '?': '|question_mark|',
            '(': '|left_parentheses|',
            ')': '|right_parentheses|',
            '--': '|dash|',
            '\n': '|return|'
        }

        return tokenize_dict

        """
        DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
        """
        tests.test_tokenize(token_lookup)
```

Tests Passed

Preprocess all the data and save it

Running the code cell below will preprocess all the data and save it to file.

```
In [6]: """
        DON'T MODIFY ANYTHING IN THIS CELL
        """
        # Preprocess Training, Validation, and Testing Data
        helper.preprocess_and_save_data(data_dir, token_lookup, create_lookup_tables)
```

Check Point

This is your first checkpoint. If you ever decide to come back to this notebook or have to restart the notebook, you can start from here. The preprocessed data has been saved to disk.

```
In [1]: """
DON'T MODIFY ANYTHING IN THIS CELL
"""

import helper
import numpy as np
import problem_unittests as tests

int_text, vocab_to_int, int_to_vocab, token_dict = helper.load_preprocess()
```

Build the Neural Network

You'll build the components necessary to build a RNN by implementing the following functions below:

- get_inputs
- get_init_cell
- get_embed
- build_rnn
- build_nn
- get_batches

Check the Version of TensorFlow and Access to GPU

```
In [2]: """
DON'T MODIFY ANYTHING IN THIS CELL
"""

from distutils.version import LooseVersion
import warnings
import tensorflow as tf

# Check TensorFlow Version
assert LooseVersion(tf.__version__) >= LooseVersion('1.3'), 'Please use TensorFlow version 1.3 or newer'
print('TensorFlow Version: {}'.format(tf.__version__))

# Check for a GPU
if not tf.test.gpu_device_name():
    warnings.warn('No GPU found. Please use a GPU to train your neural network.')
else:
    print('Default GPU Device: {}'.format(tf.test.gpu_device_name()))

TensorFlow Version: 1.3.0
Default GPU Device: /gpu:0
```

Input

Implement the `get_inputs()` function to create TF Placeholders for the Neural Network. It should create the following placeholders:

- Input text placeholder named "input" using the [TF Placeholder](https://www.tensorflow.org/api_docs/python/tf/placeholder) (https://www.tensorflow.org/api_docs/python/tf/placeholder) name parameter.
- Targets placeholder
- Learning Rate placeholder

Return the placeholders in the following tuple (`Input`, `Targets`, `LearningRate`)

```
In [3]: def get_inputs():
        """
        Create TF Placeholders for input, targets, and learning rate.
        :return: Tuple (input, targets, learning rate)
        """
        # TODO: Implement Function
        inputs = tf.placeholder(tf.int32, shape=(None, None), name='input')
        targets = tf.placeholder(tf.int32, shape=(None, None), name='targets'
        )
        learning_rate = tf.placeholder(tf.float32, name='learning_rate')
        return inputs, targets, learning_rate

        """
        DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
        """
        tests.test_get_inputs(get_inputs)
```

Tests Passed

Build RNN Cell and Initialize

Stack one or more [BasicLSTMCells](https://www.tensorflow.org/api_docs/python/tf/contrib/rnn/BasicLSTMCell)

(https://www.tensorflow.org/api_docs/python/tf/contrib/rnn/BasicLSTMCell), in a [MultiRNNCell](https://www.tensorflow.org/api_docs/python/tf/contrib/rnn/MultiRNNCell) (https://www.tensorflow.org/api_docs/python/tf/contrib/rnn/MultiRNNCell).

- The Rnn size should be set using `rnn_size`
- Initialize Cell State using the MultiRNNCell's [zero_state\(\)](https://www.tensorflow.org/api_docs/python/tf/contrib/rnn/MultiRNNCell#zero_state). (https://www.tensorflow.org/api_docs/python/tf/contrib/rnn/MultiRNNCell#zero_state) function
 - Apply the name "initial_state" to the initial state using [tf.identity\(\)](https://www.tensorflow.org/api_docs/python/tf/identity). (https://www.tensorflow.org/api_docs/python/tf/identity).

Return the cell and initial state in the following tuple (`Cell`, `InitialState`)

```

In [4]: def add_dropout(lstm_cell, keep_prob):
        ''' Add dropout to the cell.

        Arguments
        -----
        lstm_cell: tensorflow BasicLSTMCell() object
        keep_prob: Scalar tensor (tf.placeholder) for the
                   dropout keep probability

        '''
        return tf.contrib.rnn.DropoutWrapper(lstm_cell, output_keep_prob=keep_prob)

def get_init_cell(batch_size, rnn_size):
    """
    Create an RNN Cell and initialize it.
    :param batch_size: Size of batches
    :param rnn_size: Size of RNNs
    :return: Tuple (cell, initialize state)
    """
    # TODO: Implement Function

    # Probability an LSTM cell will not be dropped.
    keep_prob = 0.5
    # Create an LSTM cell
    cell = tf.contrib.rnn.BasicLSTMCell(rnn_size)

    # Stack multiple LSTM cells according to the number
    # of layers (rnn_layers). Wrap each LSTM cell in a
    # dropout layer.
    cell = tf.nn.rnn_cell.MultiRNNCell([add_dropout(cell, keep_prob)])

    # Create initial state.
    initial_state = tf.identity(cell.zero_state(batch_size, tf.float32),
                                name='initial_state')
    return cell, initial_state

"""
DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
"""
tests.test_get_init_cell(get_init_cell)

```

Tests Passed

Word Embedding

Apply embedding to input_data using TensorFlow. Return the embedded sequence.


```
In [5]: def get_embed(input_data, vocab_size, embed_dim):
        """
        Create embedding for <input_data>.
        :param input_data: TF placeholder for text input.
        :param vocab_size: Number of words in vocabulary.
        :param embed_dim: Number of embedding dimensions
        :return: Embedded input.
        """

        # TODO: Implement Function
        embedding = tf.Variable(tf.random_uniform(shape=(vocab_size, embed_d
im),
                                                minval=-1, maxval=1))

        embed = tf.nn.embedding_lookup(params=embedding, ids=input_data)
        return embed

        """
        DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
        """
        tests.test_get_embed(get_embed)
```

Tests Passed

Build RNN

You created a RNN Cell in the `get_init_cell()` function. Time to use the cell to create a RNN.

- Build the RNN using the `tf.nn.dynamic_rnn()`.
(https://www.tensorflow.org/api_docs/python/tf/nn/dynamic_rnn)
 - Apply the name "final_state" to the final state using `tf.identity()`.
(https://www.tensorflow.org/api_docs/python/tf/identity)

Return the outputs and final_state state in the following tuple (Outputs, FinalState)

```
In [6]: def build_rnn(cell, inputs):
        """
        Create a RNN using a RNN Cell
        :param cell: RNN Cell
        :param inputs: Input text data
        :return: Tuple (Outputs, Final State)
        """

        # TODO: Implement Function
        outputs, final_state = tf.nn.dynamic_rnn(cell, inputs, dtype=tf.float32)

        final_state = tf.identity(final_state, name='final_state')

        return outputs, final_state

        """
        DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
        """
        tests.test_build_rnn(build_rnn)
```

Tests Passed

Build the Neural Network

Apply the functions you implemented above to:

- Apply embedding to input_data using your get_embed(input_data, vocab_size, embed_dim) function.
- Build RNN using cell and your build_rnn(cell, inputs) function.
- Apply a fully connected layer with a linear activation and vocab_size as the number of outputs.

Return the logits and final state in the following tuple (Logits, FinalState)

```

In [7]: def build_nn(cell, rnn_size, input_data, vocab_size, embed_dim):
        """
        Build part of the neural network
        :param cell: RNN cell
        :param rnn_size: Size of rnns
        :param input_data: Input data
        :param vocab_size: Vocabulary size
        :param embed_dim: Number of embedding dimensions
        :return: Tuple (Logits, FinalState)
        """
        # TODO: Implement Function
        embedding_layer = get_embed(input_data, vocab_size, embed_dim)
        rnn_outputs, final_state = build_rnn(cell, embedding_layer)
        logits = tf.contrib.layers.fully_connected(inputs=rnn_outputs, num_o
utputs=vocab_size,
                                                    activation_fn=None)

        return logits, final_state

        """
        DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
        """
        tests.test_build_nn(build_nn)

```

Tests Passed

Batches

Implement `get_batches` to create batches of input and targets using `int_text`. The batches should be a Numpy array with the shape (number of batches, 2, batch size, sequence length). Each batch contains two elements:

- The first element is a single batch of **input** with the shape [batch size, sequence length]
- The second element is a single batch of **targets** with the shape [batch size, sequence length]

If you can't fill the last batch with enough data, drop the last batch.

For example, `get_batches([1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20], 3, 2)` would return a Numpy array of the following:

```
[
  # First Batch
  [
    # Batch of Input
    [[ 1  2], [ 7  8], [13 14]]
    # Batch of targets
    [[ 2  3], [ 8  9], [14 15]]
  ]

  # Second Batch
  [
    # Batch of Input
    [[ 3  4], [ 9 10], [15 16]]
    # Batch of targets
    [[ 4  5], [10 11], [16 17]]
  ]

  # Third Batch
  [
    # Batch of Input
    [[ 5  6], [11 12], [17 18]]
    # Batch of targets
    [[ 6  7], [12 13], [18  1]]
  ]
]
```

Notice that the last target value in the last batch is the first input value of the first batch. In this case, 1. This is a common technique used when creating sequence batches, although it is rather unintuitive.


```

In [8]: def get_batches(int_text, batch_size, seq_length):
        """
        Return batches of input and target
        :param int_text: Text with the words replaced by their ids
        :param batch_size: The size of batch
        :param seq_length: The length of sequence
        :return: Batches as a Numpy array
        """

        # TODO: Implement Function

        # Keep only enough words in int_text to make full batches.
        words_per_batch = batch_size * seq_length
        number_of_batches = len(int_text) // words_per_batch
        words_to_keep = words_per_batch * number_of_batches
        int_text = np.array(int_text[:words_to_keep])

        # Prepare for creation of the target values:
        # Target value of an element at a given position in int_text
        # is that of the element residing at the subsequent position.
        target_int_text = int_text.copy()
        # Move the first element of the input text to the final position.
        # This handles the edge case of how to choose a target value for
        # the final element of int_text.
        target_int_text[0:-1], target_int_text[-1] = int_text[1:], int_text[
0]

        # Reshape int_text to have the same number of rows
        # as the size of the batches.
        int_text = int_text.reshape((batch_size, -1))

        # Also reshape target_int_text to have the same number of rows
        # as the size of the batches.
        target_int_text = target_int_text.reshape((batch_size, -1))

        # An array to contain all the batches
        all_batches = []

        # Slice features (x) and targets (y) columns from
        # their respective arrays (int_text, target_int_text)
        # in order to populate the batches. Each batch contains
        # one sequence of values from each row of the arrays,
        # for both features and targets.
        # The start and stop of each sequence (or the indices of
        # of the columns that are sliced) correspond to the
        # seq_length parameter's value.
        for n in range(0, int_text.shape[1], seq_length):
            x_batch = int_text[:, n:n+seq_length]
            y_batch = target_int_text[:, n:n+seq_length]
            all_batches.append([x_batch, y_batch])

        # Make sure an np.array is returned
        return np.array(all_batches)

        """
DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE

```

```
"""
tests.test_get_batches(get_batches)
```

Tests Passed

Neural Network Training

Hyperparameters

Tune the following parameters:

- Set `num_epochs` to the number of epochs.
- Set `batch_size` to the batch size.
- Set `rnn_size` to the size of the RNNs.
- Set `embed_dim` to the size of the embedding.
- Set `seq_length` to the length of sequence.
- Set `learning_rate` to the learning rate.
- Set `show_every_n_batches` to the number of batches the neural network should print progress.

```
In [52]: # Number of Epochs
num_epochs = 1000
# Batch Size
batch_size = 1000
# RNN Size
rnn_size = 512
# Embedding Dimension Size
embed_dim = 300
# Sequence Length
seq_length = 10
# Learning Rate
learning_rate = 0.001
# Show stats for every n number of batches
show_every_n_batches = 60

"""
DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
"""

save_dir = './save'
```

Build the Graph

Build the graph using the neural network you implemented.

```
In [53]: """
DON'T MODIFY ANYTHING IN THIS CELL
"""

from tensorflow.contrib import seq2seq

train_graph = tf.Graph()
with train_graph.as_default():
    vocab_size = len(int_to_vocab)
    input_text, targets, lr = get_inputs()
    input_data_shape = tf.shape(input_text)
    cell, initial_state = get_init_cell(input_data_shape[0], rnn_size)
    logits, final_state = build_nn(cell, rnn_size, input_text, vocab_size,
    embed_dim)

    # Probabilities for generating words
    probs = tf.nn.softmax(logits, name='probs')

    # Loss function
    cost = seq2seq.sequence_loss(
        logits,
        targets,
        tf.ones([input_data_shape[0], input_data_shape[1]]))

    # Optimizer
    optimizer = tf.train.AdamOptimizer(lr)

    # Gradient Clipping
    gradients = optimizer.compute_gradients(cost)
    capped_gradients = [(tf.clip_by_value(grad, -1., 1.), var) for grad,
    var in gradients if grad is not None]
    train_op = optimizer.apply_gradients(capped_gradients)
```

Train

Train the neural network on the preprocessed data. If you have a hard time getting a good loss, check the [forums \(https://discussions.udacity.com/\)](https://discussions.udacity.com/) to see if anyone is having the same problem.


```

In [54]: """
DON'T MODIFY ANYTHING IN THIS CELL
"""

batches = get_batches(int_text, batch_size, seq_length)

with tf.Session(graph=train_graph) as sess:
    sess.run(tf.global_variables_initializer())

    for epoch_i in range(1, num_epochs+1):
        state = sess.run(initial_state, {input_text: batches[0][0]})

        for batch_i, (x, y) in enumerate(batches):
            feed = {
                input_text: x,
                targets: y,
                initial_state: state,
                lr: learning_rate}
            train_loss, state, _ = sess.run([cost, final_state, train_op], feed)

            # Show every <show_every_n_batches> batches
            if ((epoch_i-1) * len(batches) + batch_i+1) % show_every_n_batches == 0:
                print('Epoch {:>3} Batch {:>4}/{}    train_loss = {:.3f}'
                    .format(
                        epoch_i,
                        batch_i+1,
                        len(batches),
                        train_loss))

        # Save Model
        saver = tf.train.Saver()
        saver.save(sess, save_dir)
        print('Model Trained and Saved')

```

Epoch	10	Batch	6/6	train_loss = 5.506
Epoch	20	Batch	6/6	train_loss = 4.901
Epoch	30	Batch	6/6	train_loss = 4.466
Epoch	40	Batch	6/6	train_loss = 4.127
Epoch	50	Batch	6/6	train_loss = 3.835
Epoch	60	Batch	6/6	train_loss = 3.574
Epoch	70	Batch	6/6	train_loss = 3.344
Epoch	80	Batch	6/6	train_loss = 3.125
Epoch	90	Batch	6/6	train_loss = 2.937
Epoch	100	Batch	6/6	train_loss = 2.755
Epoch	110	Batch	6/6	train_loss = 2.597
Epoch	120	Batch	6/6	train_loss = 2.465
Epoch	130	Batch	6/6	train_loss = 2.310
Epoch	140	Batch	6/6	train_loss = 2.198
Epoch	150	Batch	6/6	train_loss = 2.082
Epoch	160	Batch	6/6	train_loss = 1.978
Epoch	170	Batch	6/6	train_loss = 1.855
Epoch	180	Batch	6/6	train_loss = 1.780
Epoch	190	Batch	6/6	train_loss = 1.657
Epoch	200	Batch	6/6	train_loss = 1.582
Epoch	210	Batch	6/6	train_loss = 1.505
Epoch	220	Batch	6/6	train_loss = 1.442
Epoch	230	Batch	6/6	train_loss = 1.382
Epoch	240	Batch	6/6	train_loss = 1.283
Epoch	250	Batch	6/6	train_loss = 1.231
Epoch	260	Batch	6/6	train_loss = 1.167
Epoch	270	Batch	6/6	train_loss = 1.121
Epoch	280	Batch	6/6	train_loss = 1.069
Epoch	290	Batch	6/6	train_loss = 1.013
Epoch	300	Batch	6/6	train_loss = 0.965
Epoch	310	Batch	6/6	train_loss = 0.925
Epoch	320	Batch	6/6	train_loss = 0.893
Epoch	330	Batch	6/6	train_loss = 0.869
Epoch	340	Batch	6/6	train_loss = 0.808
Epoch	350	Batch	6/6	train_loss = 0.787
Epoch	360	Batch	6/6	train_loss = 0.765
Epoch	370	Batch	6/6	train_loss = 0.741
Epoch	380	Batch	6/6	train_loss = 0.710
Epoch	390	Batch	6/6	train_loss = 0.684
Epoch	400	Batch	6/6	train_loss = 0.668
Epoch	410	Batch	6/6	train_loss = 0.654
Epoch	420	Batch	6/6	train_loss = 0.649
Epoch	430	Batch	6/6	train_loss = 0.623
Epoch	440	Batch	6/6	train_loss = 0.595
Epoch	450	Batch	6/6	train_loss = 0.591
Epoch	460	Batch	6/6	train_loss = 0.579
Epoch	470	Batch	6/6	train_loss = 0.563
Epoch	480	Batch	6/6	train_loss = 0.561
Epoch	490	Batch	6/6	train_loss = 0.554
Epoch	500	Batch	6/6	train_loss = 0.545
Epoch	510	Batch	6/6	train_loss = 0.529
Epoch	520	Batch	6/6	train_loss = 0.525
Epoch	530	Batch	6/6	train_loss = 0.512
Epoch	540	Batch	6/6	train_loss = 0.505
Epoch	550	Batch	6/6	train_loss = 0.508
Epoch	560	Batch	6/6	train_loss = 0.503
Epoch	570	Batch	6/6	train_loss = 0.493

```
Epoch 580 Batch      6/6   train_loss = 0.491
Epoch 590 Batch      6/6   train_loss = 0.483
Epoch 600 Batch      6/6   train_loss = 0.481
Epoch 610 Batch      6/6   train_loss = 0.467
Epoch 620 Batch      6/6   train_loss = 0.467
Epoch 630 Batch      6/6   train_loss = 0.465
Epoch 640 Batch      6/6   train_loss = 0.466
Epoch 650 Batch      6/6   train_loss = 0.456
Epoch 660 Batch      6/6   train_loss = 0.460
Epoch 670 Batch      6/6   train_loss = 0.451
Epoch 680 Batch      6/6   train_loss = 0.444
Epoch 690 Batch      6/6   train_loss = 0.444
Epoch 700 Batch      6/6   train_loss = 0.444
Epoch 710 Batch      6/6   train_loss = 0.432
Epoch 720 Batch      6/6   train_loss = 0.442
Epoch 730 Batch      6/6   train_loss = 0.432
Epoch 740 Batch      6/6   train_loss = 0.444
Epoch 750 Batch      6/6   train_loss = 0.436
Epoch 760 Batch      6/6   train_loss = 0.433
Epoch 770 Batch      6/6   train_loss = 0.426
Epoch 780 Batch      6/6   train_loss = 0.426
Epoch 790 Batch      6/6   train_loss = 0.418
Epoch 800 Batch      6/6   train_loss = 0.423
Epoch 810 Batch      6/6   train_loss = 0.424
Epoch 820 Batch      6/6   train_loss = 0.422
Epoch 830 Batch      6/6   train_loss = 0.417
Epoch 840 Batch      6/6   train_loss = 0.417
Epoch 850 Batch      6/6   train_loss = 0.413
Epoch 860 Batch      6/6   train_loss = 0.414
Epoch 870 Batch      6/6   train_loss = 0.416
Epoch 880 Batch      6/6   train_loss = 0.405
Epoch 890 Batch      6/6   train_loss = 0.409
Epoch 900 Batch      6/6   train_loss = 0.410
Epoch 910 Batch      6/6   train_loss = 0.407
Epoch 920 Batch      6/6   train_loss = 0.402
Epoch 930 Batch      6/6   train_loss = 0.403
Epoch 940 Batch      6/6   train_loss = 0.402
Epoch 950 Batch      6/6   train_loss = 0.399
Epoch 960 Batch      6/6   train_loss = 0.404
Epoch 970 Batch      6/6   train_loss = 0.396
Epoch 980 Batch      6/6   train_loss = 0.399
Epoch 990 Batch      6/6   train_loss = 0.396
Epoch 1000 Batch     6/6   train_loss = 0.391
Model Trained and Saved
```

Save Parameters

Save `seq_length` and `save_dir` for generating a new TV script.

```
In [55]: """
DON'T MODIFY ANYTHING IN THIS CELL
"""

# Save parameters for checkpoint
helper.save_params((seq_length, save_dir))
```

Checkpoint

```
In [56]: """
DON'T MODIFY ANYTHING IN THIS CELL
"""

import tensorflow as tf
import numpy as np
import helper
import problem_unittests as tests

_, vocab_to_int, int_to_vocab, token_dict = helper.load_preprocess()
seq_length, load_dir = helper.load_params()
```

Implement Generate Functions

Get Tensors

Get tensors from `loaded_graph` using the function `get_tensor_by_name()`. (https://www.tensorflow.org/api_docs/python/tf/Graph#get_tensor_by_name). Get the tensors using the following names:

- "input:0"
- "initial_state:0"
- "final_state:0"
- "probs:0"

Return the tensors in the following tuple (`InputTensor`, `InitialStateTensor`, `FinalStateTensor`, `ProbsTensor`)

```
In [57]: def get_tensors(loaded_graph):
        """
        Get input, initial state, final state, and probabilities tensor from
        <loaded_graph>
        :param loaded_graph: TensorFlow graph loaded from file
        :return: Tuple (InputTensor, InitialStateTensor, FinalStateTensor, P
        robsTensor)
        """
        # TODO: Implement Function
        input_tensor = loaded_graph.get_tensor_by_name('input:0')
        initial_state_tensor = loaded_graph.get_tensor_by_name('initial_stat
        e:0')
        final_state_tensor = loaded_graph.get_tensor_by_name('final_state:0'
        )
        probs_tensor = loaded_graph.get_tensor_by_name('probs:0')
        return input_tensor, initial_state_tensor, final_state_tensor, probs
        _tensor

        """
        DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
        """
        tests.test_get_tensors(get_tensors)
```

Tests Passed

Choose Word

Implement the `pick_word()` function to select the next word using probabilities.

```
In [58]: def pick_word(probabilities, int_to_vocab):
        """
        Pick the next word in the generated text
        :param probabilities: Probabilites of the next word
        :param int_to_vocab: Dictionary of word ids as the keys and words as
        the values
        :return: String of the predicted word
        """
        # TODO: Implement Function

        # Randomly choose a word from the four words that have
        # the highest probabilities.
        top_n = 4

        # Get the indices corresponding to the top probilities
        idx_top_probs = np.argpartition(probabilities, -top_n)[-top_n:]

        # Choose one of these indices at random
        choice = np.random.choice(idx_top_probs)

        # Select the word that belongs to the randomly
        # chosen index.
        next_word = int_to_vocab[choice]

        return next_word

        """
        DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
        """
        tests.test_pick_word(pick_word)
```

Tests Passed

Generate TV Script

This will generate the TV script for you. Set `gen_length` to the length of TV script you want to generate.

```

In [59]: gen_length = 500
         # homer_simpson, moe_szyslak, or Barney_Gumble
         prime_word = 'moe_szyslak'

         """
         DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
         """

         loaded_graph = tf.Graph()
         with tf.Session(graph=loaded_graph) as sess:
             # Load saved model
             loader = tf.train.import_meta_graph(load_dir + '.meta')
             loader.restore(sess, load_dir)

             # Get Tensors from loaded model
             input_text, initial_state, final_state, probs = get_tensors(loaded_g
raph)

             # Sentences generation setup
             gen_sentences = [prime_word + ':']
             prev_state = sess.run(initial_state, {input_text: np.array([[1]])})

             # Generate sentences
             for n in range(gen_length):
                 # Dynamic Input
                 dyn_input = [[vocab_to_int[word] for word in gen_sentences[-seq_
length:]]]
                 dyn_seq_length = len(dyn_input[0])

                 # Get Prediction
                 probabilities, prev_state = sess.run(
                     [probs, final_state],
                     {input_text: dyn_input, initial_state: prev_state})

                 pred_word = pick_word(probabilities[0][dyn_seq_length-1], int_to
_vocab)

                 gen_sentences.append(pred_word)

             # Remove tokens
             tv_script = ' '.join(gen_sentences)
             for key, token in token_dict.items():
                 ending = ' ' if key in ['\n', '(', '"'] else ''
                 tv_script = tv_script.replace(' ' + token.lower(), key)
             tv_script = tv_script.replace('\n ', '\n')
             tv_script = tv_script.replace('( ', '(')

             print(tv_script)

```

INFO:tensorflow:Restoring parameters from ./save
moe_szyslak: yeah on the ball thing.
lenny_leonard: i see this for now since there who makes sense home she.
while.
carl_carlson: pick up! what?
moe_szyslak: just what have money for the president of course that old
friend, boxing i was! another moe
barney_gumble: hey! what, you're the fat one else stupid(to, lenny home
r? hey) they the should put on how a daddy? 'cause you sayin' you look,
but that you've had all a feeling moe_szyslak: him for / yeah with the"
ned_flanders: hey guys stink, usin'. it's husband at
homer_simpson:(realizing at a man sorry homer then) there, didn't be" e
ightball" problem?!
homer_simpson: thanks for homer?. uh, would.
apu_nahasapeemapetilon: i could this guy for your old man passed, two w
anted ya got in good?
barney_gumble:(a beer to the real"-- then tonight, then singing as mo
e's.. some with dirt.(counting sign at the bar)
, ah, who homer_simpson: on the phone and drink where? i! i own losers
all better again
barney_gumble: that's alright what does she got about the money
homer_simpson: and i are you! 'cause i will say no trouble. could pull
it off my face! ow? my poor man!
david_byrne: wait right the party at this bar is(a beer) what are you?,
didn't? the one, hey(determined mean a day that's drive a body? where y
ou could do up because they enveloped me with marge 'cause" joe could"
save no".
moe_szyslak: the super are my best friend is my one? you've have to be
second, booze like(with each woman).(art). i could really do it, you! a
m(raises prank watch) better were outta the the man. the only" no girl,
of any? hey that's what its fondest me last with alcohol again!. whee h
ave in there you open up? but. not to their choice. and it's

lisa_simpson: a duff!
barney_gumble: you like for be sorry, homer the clone). you know, no wa
s from here?.. it, uh., i wanna have a job for well that you still in"
here no" i wanna need? another.
seymour_skinner:(sings, disgusted lloyd to us sits there around, stick
tonight. i'll have. to uh.(takes out cell phone
bart_simpson: hi while did yourself is here puff in the car!
moe_szyslak: i mean. now i let,

The TV Script is Nonsensical

It's ok if the TV script doesn't make any sense. We trained on less than a megabyte of text. In order to get good results, you'll have to use a smaller vocabulary or get more data. Luckily there's more data! As we mentioned in the beginning of this project, this is a subset of another dataset (<https://www.kaggle.com/wcukierski/the-simpsons-by-the-data>). We didn't have you train on all the data, because that would take too long. However, you are free to train your neural network on all the data. After you complete the project, of course.

Submitting This Project

When submitting this project, make sure to run all the cells before saving the notebook. Save the notebook file as "d1nd_tv_script_generation.ipynb" and save it as a HTML file under "File" -> "Download as". Include the "helper.py" and "problem_unittests.py" files in your submission.